

STRICTLY INTERNAL - NOT FOR CLIENT DISTRIBUTION

Data Audit Playbook

Commercial guidance, client recommendation logic, risk framing, and sequenced implementation plan.

Prepared for

Mindarora LLP

Prepared by

Suhas Bhairav, Mindarora LLP

Source dataset

erp.csv

Recommended path

Custom GenAI and predictive intelligence program

MATURITY

90 / 100

AI Scale Ready

ROWS

13

records profiled

COLUMNS

8

fields mapped

Document purpose

Use this document to prepare the client conversation, prioritize paid work, and anchor recommendations in observed evidence.

CONFIDENTIALITY AND OWNERSHIP

Confidential and proprietary. This report, analysis, recommendations, workflows, and all related materials are the property of Mindarora LLP. Prepared by Suhas Bhairav. Copyright 2026 Mindarora LLP. All rights reserved.

Mindarora LLP | Suhas Bhairav

Copyright 2026 Mindarora LLP. All rights reserved. Generated 2026-05-31 UTC.

Executive Summary


erp.csv is structurally ingestible and contains a usable operating snapshot, but the enterprise readiness level is constrained by 5.8% average missingness, 1 detected quality anomalies, and 2 sensitive column classifications.

The recommended posture is 'Custom GenAI and predictive intelligence program'. The dataset currently sits in the 'AI Scale Ready' band, so transformation should sequence ontology standardization, data contracts, and workflow automation before deeper model autonomy.

Board-Level Findings

- Overall completeness is 94.2%, which determines how much of the business process can be trusted for automation.
- 0 fields require remediation before they should drive executive metrics or AI decisions.
- 1 anomalies indicate validation gaps in upstream systems or manual entry workflows.
- 2 PII-bearing fields need governance before the data can be safely exposed to assistants or copilots.
- The dataset has measures and timestamps, making it suitable for operational dashboards after cleanup.

Readiness



Advisor Steps

- 1. Open with the maturity score and position the engagement as risk reduction plus automation acceleration.
- 2. Show the client the domain model first so they see that this is more than a CSV audit.
- 3. Recommend a 2-4 week foundation sprint before promising custom GenAI outcomes.
- 4. Package low-code workflow automation as the first ROI milestone while data contracts mature.
- 5. Use the risk register to justify governance, PII controls, and source validation work.
- 6. Create a canonical entity dictionary for customers, accounts, deals, orders, activities, and financial measures.
- 7. Assign data owners for each business domain and agree minimum viable data contracts for new records.
- 8. Classify PII handling rules, retention policy, masking requirements, and access controls before broad analyst rollout.
- 9. Use timestamp fields to build lifecycle analytics and identify automation trigger points.

Options

Off-the-shelf workflow automation

100 fit

Automated reminders, enrichment, routing, validation, and manager visibility without custom model ri

Prerequisites: Stable source exports or API access, Named owners for trigger fields, Clear exception handling

Tools: Zapier, Make, n8n, Airtable Automations, HubSpot Workflows

Custom GenAI operating layer

93 fit

Role-specific assistants, semantic search, automated briefings, next-best-action recommendations, an

Prerequisites: PII policy, High-confidence ontology, Retrieval corpus or clean structured records

Tools: OpenAI API, LangGraph, LlamaIndex, vector database, warehouse semantic layer

Data foundation and ontology sprint

92 fit

Canonical entities, data contracts, source-of-truth decisions, data quality gates, and board-ready d

Prerequisites: Access to raw exports, Stakeholder interviews, System inventory

Tools: dbt, Postgres, DuckDB, BigQuery, Snowflake

Where AI Helps Realistically

Use these points to keep the client conversation grounded: AI is valuable where it explains, summarizes, routes, and assists review; it should not bypass data contracts or owner approval.

Decision support for managers and consultants

AI can summarize the audit, explain quality gaps, and turn findings into prioritized action lists.

Evidence: 0 low-completeness fields need sequencing; 1 anomaly signals need triage

Prerequisites: Canonical business entities, Data owners for exceptions, Reusable report templates

Limit: AI can recommend priorities, but owners still need to confirm operational context.

Workflow triage and exception explanation

AI can draft explanations for flagged records and group similar exceptions for operations teams.

Evidence: 2 lifecycle timestamp fields available; 2 numeric measures available

Prerequisites: Exception taxonomy, Approved message templates, Human review queue

Limit: Rules should detect exceptions; AI should explain and route them.

Workflows

Lifecycle exception monitor

High fit

Flags stuck records, stale deals, delayed orders, or overdue cases before revenue leakage compounds.

- Implementation: Rule engine triggers alerts when records remain in a stage beyond agreed service thresholds.

Financial or pipeline anomaly review

Medium fit

Routes impossible values, outliers, and suspicious changes to operations owners before reporting cycles.

- Implementation: Validation checks run on ingest and create review tasks for records outside accepted business ranges.

Customer 360 enrichment queue

Medium fit

Improves segmentation, personalization, and lead routing by filling missing firmographic or contact attributes.

- Implementation: Missing critical fields trigger enrichment, deduplication, and account-owner follow-up workflows.

System Fingerprint

- Source type: CSV / Excel flat-file export
- Likely systems: ERP / Finance
- Join candidates: id, tax_id
- Lifecycle timestamps: order_date, ship_date
- Measures: id, invoice_amount

Domain Model

Customer - 87.2% complete

Fields: company_name, email, tax_id

Invoice - 98.5% complete

Fields: id, order_date, ship_date, invoice_amount, payment_status

- Risk: invoice_amount has 1 anomalies

Risk Register

- MEDIUM - Metric Trust: 1 anomalies detected across profiled columns. Mitigation: Add validation at source, quarantine bad records, and publish exception dashboards.
- HIGH - Privacy and Access Control: 2 PII columns detected. Mitigation: Apply masking, role-based access, retention policy, and prompt-boundary rules for AI assistants.

Minimum Data Contracts

- Invoice: require ship_date, id, order_date, invoice_amount, payment_status. Assign a business data owner for Invoice and review exceptions weekly.
- Customer: require tax_id, email, company_name. Assign a business data owner for Customer and review exceptions weekly.

Agent Readiness

Lead Scoring Agent - 28% - Blocked

- Blocker: Lacks explicit lead-level attributes (industry, size, lifecycle stage) and engagement signals
- Blocker: Dataset is invoice-centric with limited historical conversion data
- Blocker: Data quality issues: duplicate records (same Order ID appears twice) and potential date inconsistencies
- Recommendation: Deduplicate records and enforce a true primary key for orders/invoices
- Recommendation: Integrate with CRM/customer master data to enrich with lead attributes (industry, company size, region, lifecycle stage)
- Recommendation: Capture or derive lead engagement signals (last contact, response rate, marketing interactions)

RAG Document Retriever - 50% - Blocked

- Blocker: No unstructured document corpus or document store connected in current context
- Blocker: Dataset contains only structured invoice data; no attached docs or doc IDs
- Recommendation: Ingest and index unstructured documents (attachments, contracts, PDFs) into a document store
- Recommendation: Create document embeddings and retrieval prompts aligned to invoice queries
- Recommendation: Link invoices to related documents (contracts, terms) to enable retrieval over relevant docs

Churn Predictor - 32% - Blocked

- Blocker: Limited longitudinal customer history; only a handful of invoice records
- Blocker: No explicit churn labels or customer tenure features
- Blocker: Quality issues: missing values (null email), inconsistent ship_date vs order_date, and duplicates
- Recommendation: Add a customer-level history dataset (monthly orders, renewals, cancellations, tenure)
- Recommendation: Engineer features such as time since last order, order frequency, average order value, payment stability
- Recommendation: Clean data: fix date inconsistencies, deduplicate, validate identifiers and contact info

Financial Estimator - 65% - Viable

- Blocker: Requires richer historical revenue data and customer-level aggregation
- Blocker: Some missing values and potential currency normalization across records
- Blocker: Duplicate records and date anomalies could bias estimates
- Recommendation: Create customer-level revenue history by aggregating invoices over time
- Recommendation: Standardize currency and validate monetary units if multi-currency is present

- Recommendation: Deduplicate and correct date anomalies; ensure data integrity before modeling

Remediation

1. HIGH - Model Constraint Violation (AI Readiness Layer)

Deployment of 'Lead Scoring Agent' is currently blocked.

- Fix the following blockers: Lacks explicit lead-level attributes (industry, size, lifecycle stage) and engagement signals; Dataset is invoice-centric with limited historical conversion data; Data quality issues: duplicate records (same Order ID appears twice) and potential date inconsistencies. Key Recommended Next Steps: Deduplicate records and enforce a true primary key for orders/invoices; Integrate with CRM/customer master data to enrich with lead attributes (industry, company size, region, lifecycle stage); Capture or derive lead engagement signals (last contact, response rate, marketing interactions); Implement robust data quality checks and lineage tracking

2. HIGH - Model Constraint Violation (AI Readiness Layer)

Deployment of 'Churn Predictor' is currently blocked.

- Fix the following blockers: Limited longitudinal customer history; only a handful of invoice records; No explicit churn labels or customer tenure features; Quality issues: missing values (null email), inconsistent ship_date vs order_date, and duplicates. Key Recommended Next Steps: Add a customer-level history dataset (monthly orders, renewals, cancellations, tenure); Engineer features such as time since last order, order frequency, average order value, payment stability; Clean data: fix date inconsistencies, deduplicate, validate identifiers and contact info; Ensure privacy by design and obtain consent for using personal data in modeling

3. MEDIUM - Data Cleansing Issue (Invoice)

Attribute 'invoice_amount' contains 1 active format/outlier anomalies.

- Review the records. Run formatting normalization script (e.g., lowercase email normalization or phone standardization) on targeted column.

4. MEDIUM - Model Constraint Violation (AI Readiness Layer)

Deployment of 'RAG Document Retriever' is currently blocked.

- Fix the following blockers: No unstructured document corpus or document store connected in current context; Dataset contains only structured invoice data; no attached docs or doc IDs. Key Recommended Next Steps: Ingest and index unstructured documents (attachments, contracts, PDFs) into a document store; Create document embeddings and retrieval prompts aligned to invoice queries; Link invoices to related documents (contracts, terms) to enable retrieval over relevant docs; Define data governance for document retention and privacy

Column Evidence Appendix

- id: int64, 100.0% complete, 0 nulls, 12 distinct, 0 anomalies
- company_name: object, 100.0% complete, 0 nulls, 13 distinct, 0 anomalies
- email: object, 84.62% complete, 2 nulls, 11 distinct, 0 anomalies

- order_date: object, 100.0% complete, 0 nulls, 12 distinct, 0 anomalies
- ship_date: object, 92.31% complete, 1 nulls, 11 distinct, 0 anomalies
- invoice_amount: float64, 100.0% complete, 0 nulls, 12 distinct, 1 anomalies
- payment_status: object, 100.0% complete, 0 nulls, 3 distinct, 0 anomalies
- tax_id: object, 76.92% complete, 3 nulls, 9 distinct, 0 anomalies